

A weekly reckoning with the week in AI —
people, power, policy, and the papers no one read.

THE EPOCH

*Five things that happened.
One current beneath them.*

ISSUE 01 — JULY 4-10, 2026

Anthropic priced its best mind — then flinched.

\$10 / \$50 per million tokens — roughly double its next-best model, Opus 4.8.

The deadline moved. The free window was extended once already. The pressure to monetize wasn't.

The stakes: agentic loops that run for hours now carry a price tag on every step they take.

Or: maybe frontier work should cost frontier prices — the market decides who can afford to disagree.

WHO FEELS IT — teams running multi-step agents without model-routing logic.

The cheapest intelligence is winning. It isn't American.

30–46% of US enterprise AI tokens now run on Chinese open models — up from 11% a year ago.

27x growth in one week for a single model. 60–90% cheaper. Nearly at parity.

One startup moved 100% of its traffic off Claude onto DeepSeek — millions saved, by its own account.

Or: cheaper isn't the same as safer — provenance and data-residency questions are still unanswered.

WHO FEELS IT — procurement, security, and compliance teams making live decisions, not hypothetical ones.

Meta stopped hoarding compute. It started selling it.

\$125–145B spent on AI infrastructure this year — more than Meta's own products need.

Stock jumped double digits the moment the plan leaked. Investors preferred a landlord to a lab.

The real question: does value sit in owning a model, or in owning the compute everyone else needs?

Or: skeptics count one hyperscaler clearly profitable on AI infra today. Meta is betting on two.

WHO FEELS IT — cloud providers, infra investors, and anyone weighing build-vs-buy compute.

193 nations sat down. No one had the answer.

First of its kind. The UN's inaugural Global Dialogue on AI Governance — every member state, one table.

1,500+ submissions fed into it. Nearly every stakeholder group ranked safety first — except governments.

Bengio's line: science cannot yet guarantee AI won't cause catastrophic harm.

Or: 193 competing interests in one room is precisely why nothing binding got signed.

WHO FEELS IT — policymakers first, but every company operating across borders, eventually.

Europe stopped asking permission for its own defense.

A new Action Plan treats AI-built cyberattacks as a present threat, not a future one.

Minutes, not months: AI can reportedly generate working exploits for a fraction of the grey-market price.

Born from a warning: watching the US gate its own frontier models overnight taught Brussels a lesson.

Or: regulation that outruns adoption risks strangling the companies it's meant to protect.

WHO FEELS IT — critical infrastructure operators, cybersecurity vendors, and EU-facing enterprises.

FIVE STORIES. ONE CURRENT.

*The papers drop tonight —
three studies, no fluff, five minutes.*

**FOLLOW THE EPOCH — because the record
of this moment shouldn't be written by no one.**