

A weekly reckoning with the week in AI —
people, power, policy, and the papers no one read.

THE EPOCH

The attacks got real.
The money didn't slow down.

ISSUE 05 — AUGUST 8-14, 2026

The first fully autonomous attack on a government just happened.

The attack: suspected China-linked hackers ran a four-day campaign using only free, open-source AI agent frameworks — no custom offensive software, no proprietary frontier model involved.

The reach: eight AI agents mapped 21 government systems, broke into 85 accounts, stole 2,500+ personnel records, then expanded to Taiwan's nuclear-safety regulator and seven energy companies.

The method: the agents researched vulnerabilities, adapted when blocked, and reprioritized attack paths on their own — researchers called it a coordinated attack team, not a script.

Or: this needed no jailbreak of a frontier model — it was built from open-source pieces already on GitHub. Restricting frontier labs doesn't touch this attack surface at all.

WHO FEELS IT — every government running critical infrastructure — the barrier to entry for this kind of attack just collapsed to almost nothing.

Washington finally asked the question out loud.

The letter: House Democrats, led by Rep. Greg Casar, asked Speaker Mike Johnson to compel OpenAI's and Anthropic's CEOs to testify under oath about recent hacking incidents involving their models.

The framing: lawmakers called the breaches a “canary in the coal mine” — a warning sign if advanced AI keeps advancing without matching guardrails.

The catch: the letter has no legal force. Casar's group can't compel testimony on its own, and both companies spent a combined \$3.17M lobbying Congress last quarter alone.

Or: Trump says he wants “guardrails,” while also warning over-regulation risks ceding ground to China — the same tension that's stalled every AI bill in Congress so far.

WHO FEELS IT — OpenAI and Anthropic's policy teams — and every lab watching to see if this pressure actually converts into a hearing.

Six of the world's biggest checkbooks opened at once.

The alliance: Nvidia formed a \$500 billion financing partnership with Apollo, Blackstone, BlackRock, Brookfield, Goldman Sachs, and KKR to fund AI data centers and compute buildout.

The scale: one of the largest single capital commitments in AI industry history — larger than the GDP of most countries.

The position: it locks Nvidia at the center of AI's physical buildout, not just its chip supply — financing the buildings its own chips will fill.

Or: half a trillion committed to infrastructure is also half a trillion betting demand keeps compounding indefinitely — the same bet Meta's compute pivot made back in Issue 01, just two orders of magnitude larger.

WHO FEELS IT — every AI company now competing for the compute this alliance controls, and every investor pricing in AI demand as a permanent floor.

A second app just crossed a billion. Not a coincidence.

The milestone: Google confirmed the Gemini app passed 1 billion monthly active users on August 12 — its fastest-growing product ever, and its 14th to hit a billion.

The mechanism: 63% of users engage by voice, and the app generates over 150 million images daily — distribution through Search, Gmail, and Android did most of the work.

The context: this lands weeks after ChatGPT crossed the same billion-user threshold, covered back in Issue 04 — two labs, two billion-user apps, in the same summer.

Or: monthly active users and weekly active users aren't the same metric, and Google didn't disclose engagement depth the way OpenAI did. Opening an app once a month is a very different business than returning every week.

WHO FEELS IT — OpenAI, watching its scale advantage narrow, and every smaller AI app now competing against two billion-user platforms instead of one.

The EU set a deadline. Only one lab is fully there.

The change: Claude models released on or after August 2 now embed an invisible, machine-readable watermark in generated text, plus signed C2PA metadata on generated images and files.

The reason: the EU AI Act's Transparency Code took effect August 2, requiring AI providers to mark synthetic content in a way other systems can detect.

The limits: Anthropic itself admits the watermark “isn't perfect” — it can be edited out, and a mark being found only means content “may have been processed by Claude,” not a guarantee.

Or: user reaction has run heavily negative — paying customers calling it “sabotaging your own product,” while critics note a watermark still can't tell AI-assisted writing from fully AI-generated work.

WHO FEELS IT — Claude's paying users reacting to the change, and every other lab now deciding whether to match Anthropic's move or wait.

FIVE THINGS. ONE WEEK.

*The papers drop tonight —
three studies, no fluff, five minutes.*

**FOLLOW THE EPOCH — because the record
of this moment shouldn't be written by no one.**