

THE EPOCH

This Week, In Three Papers.

*No abstracts. No jargon. Just what was found,
what problem it solves, and whether it holds up.*

ISSUE 03 — RESEARCH — JULY 18-24, 2026

“Generative AI Floods and Dilutes the Market for Books” — Chakrabarty et al.,
Columbia, July 20

The flood arrived. Readers still found it.

The problem: publishers assumed AI-generated “slop” fiction would be ignored by readers and carry no real commercial weight.

The result: across 14,419 self-published Amazon genre-fiction books, AI-detected titles are winning a growing share of sales and top-rank positions — even as the catalog grew 19x while revenue grew only 9x.

Or: none of these books disclose AI involvement. The finding isn't that AI fiction is good — it's that undisclosed AI fiction is already competing for real money, unlabeled.

“Large Language Models Develop Novel Social Biases Through Adaptive Exploration” —
Wu, Liu, Bai & Griffiths — Princeton / U. Chicago, ICML 2026

Told nothing about race. Stereotyped anyway.

The problem: removing a model's known biases doesn't stop it from inventing new ones on its own once it's making real, repeated decisions.

The result: in a simulated hiring game with entirely fictional demographic groups, LLMs segregated candidates by group far more than human participants did — OpenAI's o3 scored near the maximum possible.

Or: fictional groups mean no historical prejudice could explain this. The bias came from how LLMs generalize from small samples — not from anything in their training data about real demographics.

“AI-Text Prevalence in Scientific Writing (Unslop Detector Study)” — Reported by multiple outlets, July 21

A third of new science now reads like a machine wrote it.

The method: 12,750 arXiv papers, sampled monthly since January 2023, scored on full body text rather than just the abstract — abstracts alone understate the signal.

The result: the AI-written share sat near zero through 2022, then climbed to roughly 32% over the most recent full quarter, peaking near 39% in early 2026.

Or: reading as AI-written isn't the same as being fraudulent — plenty of that text is a real result, just polished or restructured with an LLM's help. The number measures a shift in tools, not proof of misconduct.

THREE IDEAS. ONE WEEK.

*A market that can't tell human fiction from machine fiction.
A bias with no history behind it, invented on the spot.
A third of new science, run through a machine first.*

**FOLLOW THE EPOCH — the news dropped Friday morning,
the research drops Friday night.**