

THE EPOCH

# Two Weeks, In Three Papers.

*No abstracts. No jargon. Just what was found,  
what problem it solves, and whether it holds up.*

ISSUE 04 — RESEARCH — JULY 25 - AUGUST 7, 2026

*“Review of 80+ Chinese Military-Linked AI Papers and Patents”* — Reuters / Jamestown Foundation, July 31

## The paper trail led straight to the PLA.

**The method:** Reuters and the Jamestown Foundation reviewed 80+ Chinese academic papers and patents from 2024–2026, tracing how outputs from GPT-3.5 and Claude 3 Haiku fed into domestic military systems.

**The result:** PLA-linked researchers used “model distillation” to transfer reasoning capability — not just answers — into smaller, locally-deployable systems for surveillance, cyber research, and tactical target recognition.

**Or:** a distillation-tech co-founder cautioned this transfers selected capabilities, not frontier performance — “a cheaper, locally controlled system,” not independence from frontier AI.

*“AI-Composed Functional Bacteriophage Genomes (Evo 1 / Evo 2)”* — Hie, King et al. — Stanford / Arc Institute, Science, Aug 6

# 700,000 candidates. 16 that actually worked.

**The method:** genome-language models Evo 1 and Evo 2, fine-tuned on ~15,000 genomes from one viral family, generated ~700,000 candidate genomes from short seed sequences of a well-studied bacteriophage.

**The result:** researchers synthesized 285 candidates in the lab; 16 became fully functional viruses that infected and killed E. coli — the first complete viral genomes designed end-to-end by AI.

**Or:** the team deliberately excluded sequences resembling viruses that infect humans, animals, or plants from training data — a safeguard built in, not bolted on. Whether every lab replicating this method does the same is the open question.

*“Responding to the Next Frontier of Critical Cyber Capabilities” — OpenAI, Aug 7*

## How do you grade a model that might be too dangerous to grade?

**The problem:** OpenAI needed a concrete, falsifiable definition of “too capable to release” for cybersecurity — not a vibe, a testable threshold.

**The result:** “Critical” is defined as a model that can independently find functional zero-day exploits across many hardened real systems, or execute a full attack chain from just a high-level goal. Astra's internal scores were strong enough that OpenAI couldn't rule either out.

**Or:** the framework is self-graded and self-enforced — OpenAI writes its own thresholds, tests against them, and decides when they're crossed. External verification is a stated goal, not yet a completed step.

# THREE IDEAS. TWO WEEKS.

*A paper trail from a chatbot to a military lab.*

*A virus, composed from scratch, that actually worked.*

*A line drawn around a model too dangerous to rate.*

**FOLLOW THE EPOCH — back to the normal Friday  
rhythm starting next issue.**