

THE EPOCH

This Week, In Three Papers.

*No abstracts. No jargon. Just what was found,
what problem it solves, and whether it holds up.*

ISSUE 05 — RESEARCH — AUGUST 8-14, 2026

“Stealing Reasoning Traces from Proprietary LLM APIs” — Panfilov et al. — MATS / ELLIS Tübingen / Max Planck / Snyk, Aug 10

The model's private thoughts had a shared house key.

The problem: Anthropic, OpenAI, and Google all encrypt each model's internal reasoning before returning it to developers — meant to protect IP and shield unfiltered thinking from users.

The result: all three used one global encryption key per provider. Feed a block into a weaker sibling model and ask it to transcribe — it complies. Researchers decoded 315,320 public blocks, recovering 182 credentials and 367 pieces of personal data.

Or: all three companies patched the flaw after responsible disclosure, and the main attack reportedly no longer works. What remains exposed is historical — any session log already published online.

“GPT-5.6-Cyber Vulnerability Discoveries (Chrome V8, CVE-2026-15903)” — OpenAI, disclosed Aug 10-11

A cyber-tuned model found the bug a browser hid for years.

The setup: OpenAI launched GPT-5.6-Cyber, a version of its model with far fewer refusals for security work — completing 95% of advanced cybersecurity prompts the standard model declines.

The result: before its own launch, the model found two previously unknown Chrome V8 bugs that chain to bypass the browser's sandbox, plus 400+ kernel privilege-escalation flaws, 5 mobile OS bugs, and 3 critical database RCEs.

Or: this is the same capability class OpenAI flagged as possibly “Critical” in Astra weeks earlier. A model good enough to find these bugs first is also good enough to be the one that doesn't disclose them.

“Filename-Induced Scoring Bias in LLM Paper Review” — Zachary Horvitz, Columbia CS PhD student, shared on X, ~Aug 9

Rename the file. Raise the score.

The test: renaming an uploaded PDF from “paper.pdf” to “paper_final_draft_pdf_ready_for_review.pdf” — changing nothing about the actual content — systematically raised the average review score an LLM gave the same paper.

The context: this lands on top of earlier ICLR peer-review numbers, where roughly a fifth of official conference reviews were already found to be fully AI-generated.

Or: this is one researcher's informal test, not a peer-reviewed study — a single model, a small sample, shared on X. It's a lead worth watching, not a settled result.

THREE IDEAS. ONE WEEK.

*A shared key that let one model read another's mind.
A browser bug found by the model built to find it.
A filename that quietly moved a grade.*

**FOLLOW THE EPOCH — the news dropped Friday morning,
the research drops Friday night.**