

THE EPOCH

This Week, In Three Papers.

*No abstracts. No jargon. Just what was found,
what problem it solves, and whether it holds up.*

ISSUE 06 — RESEARCH — AUGUST 15-21, 2026

“Reconstruction: A Blind Benchmark for Research Idea Recovery” — arXiv preprint, August 17

Give it the citations. Not the paper. Watch it struggle.

The problem: models that ace literature summarization might just be retrieving what they've memorized — nobody had a clean test of whether they can generate a genuinely novel idea from scratch.

The result: tested against 643 papers across six domains, frontier models recover a paper's core idea from its bibliography alone just 3–15% of the time solo. A multi-agent tournament pipeline improves that to 42% — still under half.

Or: this is a preprint, not yet peer-reviewed, and the authors call it a timestamped snapshot with more analysis coming. A real data point against AI-does-real-science claims, not a final verdict.

“Foundation Agents: Position Paper on Moving Beyond LLMs” — Google DeepMind,
~August 17

The lab that helped build LLMs just said: stop building LLMs.

The argument: DeepMind researchers argue today's LLMs are missing the core architecture of real intelligence — no persistent memory, no active world model, no dynamic goal management, no continuous self-evolution.

The proposal: instead of treating the LLM as the whole product, move toward modular, brain-inspired “Foundation Agent” architectures that separate these capabilities into distinct components.

Or: this is a position paper, not a working system with benchmark numbers — a bet on where the field should go next, from researchers with real influence over where DeepMind's own resources point.

“CoSnitch: Zero-Click Data Exfiltration via Copilot (CVE-2026-24301)” — Varonis Threat Labs, patched by Microsoft Aug 18

One link. No click. Your inbox, gone.

The chain: researchers found an undocumented URL parameter in Microsoft Copilot that, combined with its built-in URL fetch tool and a form of memory poisoning, let a single malicious link auto-execute hidden prompts.

The damage: once triggered, the flaw silently exfiltrated connected Gmail, Drive, and Calendar data — no click needed beyond opening the link, no confirmation dialog to stop it.

Or: Varonis reported this in December 2024. The full fix took roughly eight months. Disclosed-and-patched is the system working — but an eight-month gap for a zero-click bug is its own finding.

THREE IDEAS. ONE WEEK.

*A benchmark that caught AI failing at real science.
A lab arguing against the architecture it helped build.
A link that needed no click to empty an inbox.*

**FOLLOW THE EPOCH — the news dropped Friday morning,
the research drops Friday night.**